

# Linux Containers & Docker

December 12, 2016 @ Urbino

Speakers: *Antonio Esposito, Michele Sorcinelli*

Reference teachers: *Andrea Seraghi, Alessandro Bogliolo*



[github.com/kobe25/docker-tutorial](https://github.com/kobe25/docker-tutorial)


Code & Slides here!

# Background

Causes of **high friction in bootstrapping** a project:

- obsolete documentation
- implicit tasks or dependency versions
- too coupled to a specific system
- tricky configuration


# Virtual Machines





- reproducible environment with “configuration management” tools
- great isolation
- full operating system virtualization
- resource overhead

# Containers

- native reproducible environment
- good isolation
- only application “virtualization”
- no resource overhead



# Virtual Machines vs Containers



# History of Containers

- 2006: cgroups (control groups)
- ...
- **2013: docker**
- 2014: rkt, kubernetes
- Today: Google runs ~2 billions containers every week

# Docker: just a “container runtime”

- Provides a **user-friendly interface** to Linux kernel isolation features
- Native on Linux systems, quite integrated with OS X and Windows
- How To Install: <https://www.docker.com/products/overview>




# Isolation

Every container has its own:

- **process** and **memory** namespaces
- **file system** namespace
- **IP address** and networking

# Container Images vs Containers

1. Build a container image
  - Images are **IMMUTABLE**
2. Run 1+ container(s) of this image
  - Containers are **RUNNING**



# Containers in practice!

# app.py

```
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello_world():
    return '<h1>Hello, World!</h1>'

app.run(host='0.0.0.0', port=8000)
```

# Dockerfile

```
FROM python:3.5
```

```
RUN pip install Flask==0.11.1
```



```
COPY ./ /code/
```

```
WORKDIR /code
```

```
EXPOSE 8000
```

```
CMD python app.py
```

# Dockerfile

253.3 MB	<code>python</code> <code>latest</code> <code>3</code> <code>3.5</code> <code>3.5.2</code> <span>What's this?</span> 
4.1 MB	<code>RUN pip install Flask==0.11.1</code>
7.7 kB	<code>COPY</code> <code>dir:d11f736eeba2c0974bfb394fbcd53b75f09d8744643b7c35d088d5461ad0de8b</code> <code>in /code/</code>  <code>WORKDIR /code</code>  <code>EXPOSE 8000/tcp</code>  <code>CMD ["/bin/sh" "-c" "python app.py"]</code> 
32 bytes	

# docker command line

Show all Docker commands:

```
$ docker help
```

Show container images and container status:

```
$ docker images
```

```
$ docker ps
```

# docker command line

Pull a Docker container image:

```
$ docker pull python
```

```
$ docker pull python:3.5
```

```
$ docker pull kobe25/docker-tutorial:latest
```

```
$ docker pull quay.io/coreos/etcd:latest
```



# docker command line

Build a container image:

```
$ docker build -t app-image .
```

Create and start a container:

```
$ docker run --name app app-image
```

```
$ docker ps
```

# docker command line

Manage containers:

```
$ docker stop app
```

```
$ docker start app
```

```
$ docker rm -f app
```

Other commands: `restart`, `kill`, `inspect..`

# Permissions and Volumes

- container user != host user
- default container user is root
- container root user can do privilege escalation

Good practices:

- use a non-root user
- host UID == container UID (because of file system permission on volume mounts)

# Networking

- PTP interface between host and container
- Dedicated bridges between containers (namespacing)
- DNS inside containers (not host)
  - resolving
  - ping

# Resource Constraints

Limit CPU usage:

```
$ docker run --cpu-shares=20 mysql
```

Limit memory usage:


```
$ docker run --memory=1024 mysql
```

Others: network I/O, disk I/O...

# Multi-Container Applications

# Docker Compose

- Enables **multi-container applications**
- Offers a user-friendly **configuration file**



# docker-compose.yml

```
version: '2'
services:
  db:
    image: mongo
  app:
    build: ./
    volumes: [ './code' ]
    ports: [ '127.0.0.1:8000:8000' ]
    depends_on: [ 'db' ]
```



# docker-compose command line

Build multiple container images and run a multi-container application:

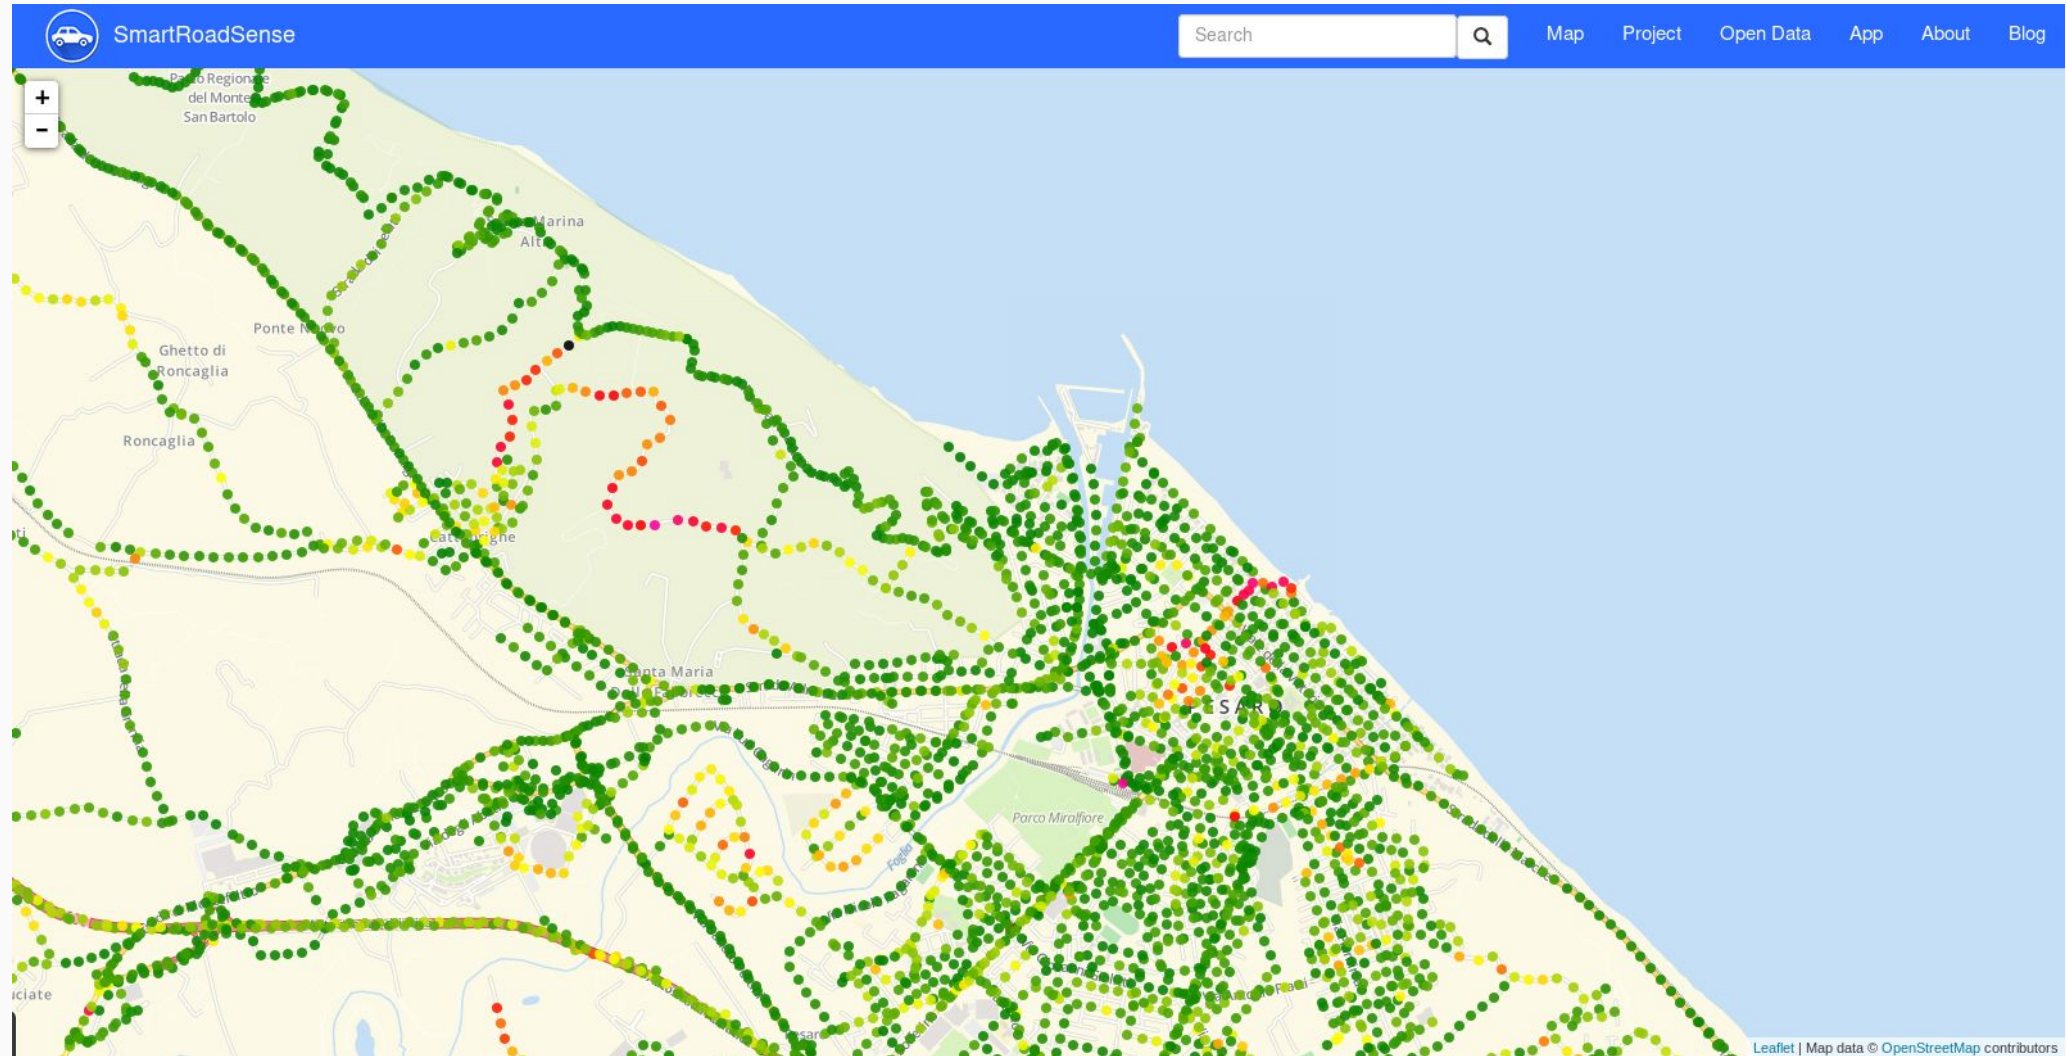
```
$ docker-compose up -d
```

Check current status:

```
$ docker-compose ps
```

# Documentation

- Docker Engine (or simply Docker): <https://docs.docker.com/engine/>
- Docker Compose: <https://docs.docker.com/compose/>



# SmartRoadSense: a real use case

## Persistent Services:

- 3 PostgreSQL/PostGIS
- 1 MySQL
- 5 PHP
- 1 Node.js
- 1 Nginx
- 1 Apache

## Scheduled Jobs:

- 2 PHP
- 1 Go
- 1 Bash

## On demand:

- 3 database clients
- 1 Node.js

# Challenges in Production

- Developers deploy 1+ application(s) to production
- High Availability requires a cluster with 3+ servers
- Optimizing resources

# Container “Production Ecosystem”

**CoreOS** is a GNU/Linux distribution optimized for containers

- low overhead, high density

**Kubernetes/OpenShift** is a Platform as a Service (PaaS)

- automating deployment, scaling, and management of containerized applications in server clusters

# Cloud Computing

**Software as a Service (SaaS)**  
Applications for End Users

**Software as a Service (SaaS)**  
Applications for End Users

**Platform as a Service (PaaS)**  
Deploying Applications by Developers

**Infrastructure as a Service (IaaS)**  
RAW Compute, Network and Storage

**Infrastructure as a Service (IaaS)**  
RAW Compute, Network and Storage

Thanks for Your Attention !

QUESTIONS?